

Artisan — Product Requirements Document

Version: 0.4

Branch: `claude/artisan-onboarding-surface-3T4Dd`

Status: In Progress

Last Updated: 2026-05-23

Table of Contents

- [Product Vision](#1-product-vision)
 - [Personas](#2-personas)
 - [Three-Stage Onboarding](#3-three-stage-onboarding)
 - [Day-in-Life Journeys](#4-day-in-life-journeys)
 - [Field-to-Matrix](#5-field-to-matrix)
 - [Behavioral Design Layer](#6-behavioral-design-layer)
 - [Stage 3 Hero Card Design](#7-stage-3-hero-card-design)
 - [Artisan as Trained Subordinate — Vagisha Trust-Building Bot](#8-artisan-as-trained-subordinate--vagisha-trust-building-bot)
 - [Data Schema](#9-data-schema)
 - [Integration Layer](#10-integration-layer)
 - [Open-Source Integration Map](#11-open-source-integration-map)
 - [Success Metrics](#12-success-metrics)
 - [Reliance Score — The Right Company Metric](#13-reliance-score--the-right-company-metric)
 - [Oracle Open Cloud POC Brief](#14-oracle-open-cloud-poc-brief)
 - [The Four Onboarding States](#15-the-four-onboarding-states)
 - [The Amit Asks Apur Quality Report](#16-the-amit-asks-apur-quality-report)
-

1. Product Vision

Artisan is an AI work surface that learns your tools, infers your role, and surfaces the right action before you know you need it. It is not a chat interface or a command-line tool — it is a daily work surface that sits between a knowledge worker and their tool stack.

Core Positioning

What Artisan Is	What Artisan Is Not
A daily work surface	A chat interface
A role-aware intelligence layer	A command-line tool
A proactive action surface	A passive assistant
A tool-stack bridge	A standalone app

YC Narrative Frame: The "First 8 Minutes" Problem

Most AI tools take too long to deliver first value. A user signs up, completes a lengthy onboarding, and still cannot point to a single concrete output that saved them time. Artisan must deliver a **specific, personalised, useful output within the onboarding session itself** — within the first 8 minutes of a user's experience.

The first 8 minutes are not setup time. They are the product.

The promise: By the time a user finishes connecting their first two tools and answers five contextual questions, Artisan surfaces one card that is immediately, demonstrably useful for that specific person on that specific day.

Design Principles

- ****Inference over configuration**** — Artisan infers role and context from tool signals; it does not ask users to fill in profiles.
 - ****Proactive over reactive**** — Artisan initiates at the right moment; it does not wait to be queried.
 - ****Evidence over assertion**** — every piece of intelligence shown carries a trust marker indicating its source.
 - ****Ritual over feature**** — the goal is to become a daily habit, not a feature used occasionally.
-

2. Personas

Artisan is designed around five primary knowledge worker personas. Each persona is defined by their tool stack, their primary pain, and the emotional hook that makes Artisan compelling to them.

2.1 PM — Meera

Attribute	Detail
Name	Meera
Role	Product Manager
Primary Tool Stack	Jira, Confluence, Slack, Google Meet, Notion
Primary Pain	Constant context-switching between sprint tracking, stakeholder updates, and async decisions. Meetings consume the day; strategic thinking happens at 9pm.
Primary Emotional Hook	"Finally, something that reads across all my tools and tells me what actually needs my attention today."
Defining Moment	Walking into a sprint planning meeting and already knowing which tickets are blocked, why, and who is the blocker — without opening Jira.

2.2 Analyst — Arjun

Attribute	Detail
Name	Arjun
Role	Business / Data Analyst
Primary Tool Stack	Google Sheets, Looker / Metabase, Slack, SQL, Notion
Primary Pain	Repetitive data pulls for ad-hoc requests that interrupt deep analysis work. The same five stakeholders ask for the same five numbers in slightly different formats.
Primary Emotional Hook	"I want to do real analysis, not be a human CSV exporter."
Defining Moment	A stakeholder asks for last week's numbers in Slack; Artisan has already drafted the reply with the current figures attached.

2.3 Marketing Manager — Vikram

Attribute	Detail
Name	Vikram
Role	Marketing Manager
Primary Tool Stack	HubSpot, Google Analytics, Canva, Slack, Notion, Mailchimp
Primary Pain	Campaign performance is scattered across five platforms. Attributing what worked requires manual aggregation every Monday morning.
Primary Emotional Hook	"Tell me what to double down on before I waste budget on another week of the wrong channel."
Defining Moment	Monday morning brief is already waiting — campaign delta, best-performing asset, recommended next action — before the weekly sync call.

2.4 Designer — Hemant / Vagisha

Attribute	Detail
Names	Hemant (IC Designer), Vagisha (Senior / Lead Designer)
Role	Product Designer
Primary Tool Stack	Figma, Notion, Slack, Loom, Linear
Primary Pain	Design review cycles are slow. Feedback is scattered across Figma comments, Slack threads, and Loom recordings. Decisions are hard to trace.
Primary Emotional Hook	"I want every reviewer to show up prepared. I'm tired of re-explaining context that's already in the file."
Defining Moment	Before a design crit, every invited reviewer has already seen the brief, the constraints, and the three questions that need decisions — surfaced by Artisan automatically.

2.5 Sales — Shubham

Attribute	Detail
Name	Shubham
Role	Account Executive / Sales
Primary Tool Stack	Salesforce / HubSpot CRM, Gmail, Google Calendar, Gong / Chorus, LinkedIn
Primary Pain	Prep time for discovery and demo calls is high. CRM hygiene is a constant guilt. Post-call follow-up degrades under pressure.
Primary Emotional Hook	"Walk me into every call knowing what the prospect cares about, what we said last time, and what the next step should be."
Defining Moment	

Attribute	Detail
	5 minutes before a demo call, a brief card surfaces with company context, last interaction summary, and a suggested opening line tailored to recent news.

3. Three-Stage Onboarding

Onboarding is not setup — it is the first proof-of-value delivery. All three stages must be completable in under 8 minutes.

Stage 1 — Connect

Purpose: Establish the data foundation. Give Artisan read access to the tools the user already uses.

What it collects:

- OAuth connections to integrated tools (Jira, Notion, Slack, Google Workspace, HubSpot, Figma, Linear, etc.)
- Metadata-only read: calendar events, ticket statuses, document titles, message threads (not full content unless explicitly granted)
- Tool presence signal — which tools exist in the stack, which are used heavily vs. lightly

What it produces:

- A ranked tool-stack map showing Artisan's coverage of the user's workflow
- An initial role inference (e.g., "Based on your tools, you appear to be a PM or engineering lead")
- A readiness score indicating how much personalisation is possible with current connections

UX Notes:

- Maximum 2 clicks per tool connection
- Show a live "coverage bar" that grows as tools are connected — immediate visual reward
- At least one tool connection must be completable in under 60 seconds

Stage 2 — Interview

Purpose: Validate and refine the role inference. Collect context that cannot be inferred from tools alone.

What it collects:

- Role confirmation or correction (select from inferred options or free-text)
- Current top 3 priorities (free-text, max 2 sentences each)
- Recurring pain: "What is the one thing you do every week that you wish you didn't have to?"
- Collaboration pattern: primarily async or meeting-heavy?
- Pre-fill from connected tools where possible (e.g., calendar density infers meeting-heavy)

What it produces:

- A persona-matched profile with confidence levels
- Pre-populated context for Stage 3 hero card generation
- A personalisation readiness assessment: "Artisan can begin briefing you on [X, Y, Z]"

UX Notes:

- Maximum 5 questions
- Every question that can be pre-answered from tool data must be pre-answered and presented for confirmation, not asked from scratch
- Questions are conversational, not form-fields

Stage 3 — Personalization

Purpose: Deliver the first concrete, personalised output. The "wow moment."

What it collects:

- User confirmation of the hero card relevance ("Is this useful? Yes / Tweak it / Not relevant")
- Feedback signal on the first output — this becomes the first training data point for the user's model

What it produces:

- One persona-specific hero card with real data from connected tools
- A suggested first daily ritual (e.g., "Check your Morning Read at 7:45am")
- An explanation of how the card was generated (trust markers, data sources)

UX Notes:

- The hero card must contain at least one data point that is specific to today and specific to this user
 - The card must be immediately actionable — it contains a primary CTA
 - The "how was this made?" explanation is available but not forced
-

4. Day-in-Life Journeys

For each persona, the journey describes the current reality ("Today's Grind"), the Artisan unlock, and the future state ("Tomorrow's Superpower"). The frame is the **full-stack builder transformation**: Artisan does not add a tool to the stack — it upgrades the entire stack simultaneously.

4.1 PM — Meera

Today's Grind

Time	Activity
8:00am	Opens Jira to check sprint status before standup
8:20am	Checks Slack for overnight messages; finds three threads needing responses
8:45am	Standup — discovers two blockers that weren't in Jira
9:30am	Confluence document for stakeholder update needs to be written from scratch
2:00pm	Back-to-back planning meetings with no prep time between
5:30pm	Finally begins the actual strategic work that was planned for 10am

The Artisan Unlock

Artisan reads across Jira, Slack, and Confluence overnight. At 7:45am, Meera's Morning Read is waiting: two blockers identified from Slack threads that haven't been filed as tickets, a draft stakeholder update pre-generated from sprint velocity data, and a 5-minute pre-meeting brief for today's planning session.

Tomorrow's Superpower

Meera walks into standup having already resolved one blocker via an async Slack nudge Artisan drafted. The stakeholder update writes itself. Strategic thinking time moves from 5:30pm to 2:00pm.

4.2 Analyst — Arjun

Today's Grind

Time	Activity
9:00am	Five ad-hoc data requests waiting in Slack from overnight
9:45am	Pulls the same weekly metrics for the third time this month

Time	Activity
11:00am	Deep analysis work begins — interrupted at 11:20am by an urgent Slack message
12:00pm	Lunch eaten at desk while completing a "quick" data pull
3:00pm	Realizes the dashboard is showing stale data; spends 45 minutes debugging
5:00pm	The analysis that was the actual goal of the day remains unfinished

The Artisan Unlock

Artisan detects recurring query patterns and pre-caches the five most-requested metrics. When a stakeholder pings for last week's numbers, Artisan surfaces a draft reply with the current figures. The "urgent" interrupt at 11:20am becomes a 30-second confirmation instead of a 40-minute task.

Tomorrow's Superpower

Arjun's deep analysis work is protected. Ad-hoc requests are handled by Artisan-drafted responses. Dashboard staleness is flagged before Arjun opens it.

4.3 Marketing Manager — Vikram

Today's Grind

Time	Activity
8:30am	Monday morning: manually pulls performance data from GA, HubSpot, Mailchimp
9:30am	Builds a slide with last week's numbers for the weekly sync
10:00am	Weekly sync — team debates which channel is working with no shared data
2:00pm	Campaign brief writing begins from scratch
4:00pm	Realizes the email campaign had a tracking link issue; can't measure correctly
5:00pm	Budget allocation decision needed by tomorrow; no clean attribution data

The Artisan Unlock

By Sunday night, Artisan has aggregated campaign performance across all channels. Monday morning brief is waiting with last week's delta, best-performing asset, anomaly flag (the tracking link issue), and a recommended budget reallocation for the week.

Tomorrow's Superpower

Vikram walks into the Monday sync with a shared brief that the whole team has already seen. Decisions are made in 20 minutes instead of 60. Budget moves faster.

4.4 Designer — Hemant / Vagisha

Today's Grind

Time	Activity
9:00am	Design crit at 10am; sends Figma link to Slack at 9:50am
10:00am	Half the reviewers haven't opened the file; 20 minutes re-explaining context
11:00am	Feedback is collected in Figma comments, Slack thread, and one Loom video
2:00pm	Tries to write up design decisions; can't find where specific feedback came from
3:30pm	Developer asks about a design decision that was discussed in a Loom nobody saved
5:00pm	Decision log is incomplete; the rationale for three choices is lost

The Artisan Unlock

Before the design crit, Artisan sends each reviewer a brief card with the Figma link, the three decisions that need input, and the constraints. During the crit, the Recording Bot captures decisions and attributions. Afterward, the ReviewTrace is auto-populated with every decision, its source, and the reviewer who made it.

Tomorrow's Superpower

Every design decision is traceable. Developers get answers in seconds. Design reviews become shorter because reviewers arrive prepared.

4.5 Sales — Shubham

Today's Grind

Time	Activity
8:00am	Three calls today; needs prep for each — opens Salesforce, Gmail, LinkedIn
8:30am	Prep for 9am call takes 25 minutes; still feels underprepared
9:00am	Discovery call — scrambling to find previous call notes mid-conversation
11:00am	Post-call: CRM update takes 15 minutes; follow-up email takes another 10
2:00pm	Demo call — prospect asks about a use case that wasn't in the prep notes
5:00pm	End of day: CRM is 2 calls behind; follow-ups are queued for tomorrow

The Artisan Unlock

At T-5 minutes before each call, Artisan surfaces a brief card: company context, last interaction summary, open questions from the previous call, and a suggested opening line based on recent news. Post-call, Artisan drafts the CRM update and follow-up email from the call transcript.

Tomorrow's Superpower

Shubham enters every call prepared. CRM hygiene is automatic. Follow-ups go out the same day. Prep time drops from 25 minutes to 5 minutes per call.

5. Field-to-Matrix

The Field-to-Matrix maps the key moments in each persona's day across eight dimensions: Time, Field Reality, Cognitive Load, Emotional State, Artisan Intervention, Emotional Resolution, Habit Hook, and Data Signal.

5.1 PM — Meera

Time	Field Reality	Cognitive Load	Emotional State	Artisan Intervention	Emotional Resolution	Habit Hook	Data Signal
7:45am	Opens phone before coffee	High — scanning 3 apps	Anxious / reactive	Morning Read card in Artisan	Relief — single surface	"Check Artisan first"	Morning open time, first-action app
8:45am	Standup — blockers surface	High — real-time synthesis	Frustrated / unprepared	Pre-standup brief with blocker flag	Confident / prepared	Brief becomes pre-standup ritual	Blocker detection lead time
12:00pm	Stakeholder update due	Very high — writing from scratch	Pressured / behind	Draft update pre-generated	Accomplished in 10 minutes	Update review becomes ritual	Doc generation acceptance rate
2:00pm	Back-to-back meetings	High — no transition time	Depleted / scattered	Pre-meeting brief at T-5	Grounded / contextual	Brief = meeting prep habit	Brief open rate, meeting prep time delta
5:15pm	End of day — what moved?	Medium	Uncertain / unresolved	EOD Close card — what moved today	Closed / resolved	EOD check becomes shutdown ritual	EOD engagement rate

5.2 Analyst — Arjun

Time	Field Reality	Cognitive Load	Emotional State	Artisan Intervention	Emotional Resolution	Habit Hook	Data Signal
9:00am	5 ad-hoc requests in Slack	High — context-switching	Overwhelmed / defensive	Artisan drafts 3 of 5 responses	Relief — protected time	Draft review replaces manual pulls	Response draft acceptance rate
11:00am	Deep analysis interrupted	Very high — broken flow	Resentful / distracted	Artisan handles "urgent" reply	Flow protected	Interrupt shield becomes trusted	Interrupt handling rate
3:00pm		High — debugging	Frustrated / wasted time		Proactive / in control	Anomaly alerts	

Time	Field Reality	Cognitive Load	Emotional State	Artisan Intervention	Emotional Resolution	Habit Hook	Data Signal
	Dashboard showing stale data			Staleness flag before manual discovery		replace manual checks	Anomaly detection accuracy
5:00pm	Core analysis unfinished	High — guilt / backlog	Defeated	EOD Close shows actual vs. planned	Calibrated / reset	EOD review becomes planning tool	Plan completion delta

5.3 Marketing Manager — Vikram

Time	Field Reality	Cognitive Load	Emotional State	Artisan Intervention	Emotional Resolution	Habit Hook	Data Signal
8:30am Monday	Manual aggregation from 4 platforms	Very high	Tedious / mechanical	Weekly brief pre-built	Energized / strategic	Monday brief replaces manual pull	Brief generation time vs. manual
10:00am	Weekly sync with no shared baseline	High — debating without data	Frustrated	Shared brief everyone has seen	Aligned / decisive	Sync prep becomes shared ritual	Sync duration delta
4:00pm	Tracking issue discovered	High — damage control	Anxious	Artisan flags anomaly + impact estimate	Contained / informed	Anomaly alerts become trusted	Anomaly flag lead time
5:00pm	Budget decision with bad attribution	Very high	Uncertain / risky	Attribution-best-available summary	Decisive despite uncertainty	Attribution summary replaces instinct	Attribution confidence score

5.4 Designer — Hemant / Vagisha

Time	Field Reality	Cognitive Load	Emotional State	Artisan Intervention	Emotional Resolution	Habit Hook	Data Signal
9:50am	Sends Figma link 10 minutes before crit	High — last-minute prep	Stressed / rushed	Pre-crit brief sent to reviewers	Confident / prepared	Brief becomes crit-prep ritual	Reviewer brief open rate
10:00am	Re-explaining context to unprepared reviewers	Very high	Frustrated / repetitive	Brief pre-answered the context	Efficient / respected	Review time drops	Crit time delta
Post-meeting	Feedback scattered	High — aggregation	Overwhelmed	ReviewTrace auto-populated	Organized / traceable	Trace review	Decision capture completeness

Time	Field Reality	Cognitive Load	Emotional State	Artisan Intervention	Emotional Resolution	Habit Hook	Data Signal
	across platforms					becomes habit	
3:30pm	Developer asks about lost Loom decision	High — archaeology	Anxious / uncertain	ReviewTrace serves the answer instantly	Reliable / authoritative	Artisan = source of truth	Decision retrieval time

5.5 Sales — Shubham

Time	Field Reality	Cognitive Load	Emotional State	Artisan Intervention	Emotional Resolution	Habit Hook	Data Signal
8:00am	Prep for 3 calls across 4 tools	Very high	Anxious / scattered	Call briefs waiting for all 3	Calm / prepared	Brief review replaces manual prep	Prep time delta per call
T-5min	Last-minute scan for context	High	Nervous / underprepared	T-5 card with opening line	Confident / sharp	T-5 card becomes pre-call ritual	T-5 card open rate
Post-call	CRM update + follow-up email	High	Drained / reluctant	CRM draft + email draft from transcript	Efficient / complete	Draft review replaces manual entry	CRM update time delta
5:00pm	2 calls behind on CRM	Medium	Guilty / behind	EOD Close — CRM health card	Clean / resolved	EOD CRM check becomes closure ritual	CRM lag time delta

6. Behavioral Design Layer

Artisan is designed to become a habit, not just a feature. The behavioral layer applies the Hooked model (Nir Eyal) to create internal triggers, investment loops, and variable rewards that make Artisan part of the user's daily rhythm.

6.1 The Hooked Model Applied

Hook Stage	Artisan Implementation
Trigger (External)	Push notification / calendar block for Morning Read; pre-meeting notification at T-5
Trigger (Internal)	"I need to know what's happening" → Artisan becomes the reflex, not Slack
Action	Open Artisan; glance at Morning Read; check pre-meeting brief
Variable Reward	Sometimes the card has a critical flag. Sometimes it reveals a pattern. Sometimes it has already resolved something.
Investment	Every confirmation ("This was useful") and correction ("Not relevant") trains the model; the product gets smarter the more you use it

6.2 The Four Rituals

Artisan is built around four named rituals. Each has a specific time window, a named card type, and a designed emotional payoff.

Morning Read — 7:45am

- ****Trigger:**** Push notification; also surfaces on first Artisan open
- ****Content:**** What moved overnight, flagged items that need attention today, one proactive action suggestion
- ****Emotional Design:**** Replace morning anxiety (scanning 3 apps) with morning clarity (one surface)
- ****Variable Reward Element:**** Occasionally surfaces an insight that isn't urgent but is strategically interesting — "did you know your top Jira contributor hasn't commented in 5 days?"

Pre-Meeting Brief — T-5 minutes

- ****Trigger:**** Calendar integration fires notification 5 minutes before any meeting with >1 attendee
- ****Content:**** Meeting context, attendee history, open items from previous related meetings, one suggested opening move
- ****Emotional Design:**** Replace "I should have prepped for this" with "I already know what I need to know"

- **Variable Reward Element:** Occasionally includes a piece of information about an attendee or topic that is unexpectedly relevant

Capture — Post-Meeting

- **Trigger:** Calendar event end + automatic recording/transcript availability
- **Content:** Draft decisions log, action items with suggested owners, follow-up message drafts
- **Emotional Design:** Replace post-meeting admin guilt with instant closure
- **Variable Reward Element:** Occasionally surfaces a decision from this meeting that contradicts a decision from 3 weeks ago — "Flag: this conflicts with the decision made in the April 12 sync"

EOD Close — 5:15pm

- **Trigger:** Time-based (configurable); also triggered by "end of day" inferred from calendar clearing
- **Content:** What moved today vs. what was planned, unresolved items carried forward, one thing to note for tomorrow morning
- **Emotional Design:** Replace "did I actually do anything today?" with structured closure
- **Variable Reward Element:** Occasionally shows a streak ("3 days in a row of clearing your blocker backlog before EOD") or a delta ("Your standup prep time dropped 40% this week")

6.3 Variable Reward Design

Variable rewards are not random. They are designed around three reward types:

Reward Type	Description	Example
Tribe	Social proof or recognition within the user's context	"Your design crit had the highest reviewer prep rate this quarter"
Hunt	Discovery of something unexpected but valuable	"Slack analysis shows a recurring blocker pattern you haven't filed as a risk"
Self	Progress, mastery, completion	"You've resolved 4 of your 5 weekly blockers before Thursday for the first time"

Each ritual is designed to deliver at least one reward of each type per week.

6.4 Trust Marker Visual Language

Every piece of intelligence shown in Artisan carries a trust marker — a visual indicator of how the information was derived. This is critical for user trust and for calibrating when to act vs. when to verify.

Marker	Color	Meaning
Verified	Green dot	Directly read from a connected tool's structured data (e.g., Jira ticket status = "Blocked")
Inferred	Amber dot	Derived by Artisan from patterns across data (e.g., "This person is likely the blocker based on thread activity")
Suggested	Blue dot	Artisan-generated recommendation with no direct data backing (e.g., "Consider following up with this stakeholder")

Trust markers appear inline with each data point, not just at the card level. Users can tap any trust marker to see its source.

6.5 The Four Scenarios Where Artisan Initiates

Artisan sends proactive notifications in exactly four scenarios — never more, never arbitrary:

Scenario	Trigger Condition	Notification Type
Ritual Window	Time-based trigger for Morning Read / EOD Close	Soft notification — "Your Morning Read is ready"
Meeting Imminent	Calendar event starting in T-5 minutes	Medium notification — "Brief for [Meeting Name] ready"
Anomaly Detected	Data signal deviates from established baseline by $>2\sigma$	High-priority flag — "Something unexpected: [summary]"
Action Required	A tracked item has been unresolved past its expected resolution window	Reminder — "[Item] has been open for [X] days — still relevant?"

Artisan does not send motivational nudges, generic tips, or unsolicited suggestions outside these four scenarios.

7. Stage 3 Hero Card Design

The hero card is the first concrete output Artisan delivers at the end of onboarding. It must feel specific, personal, and immediately useful. Generic cards are a product failure.

7.1 Card Anatomy

Every hero card contains:

```
[Role Icon] [Persona-matched headline]
[Temporal hook copy]
[Stat pill 1] [Stat pill 2] [Stat pill 3]
[Primary insight – 2 sentences max]
[Trust source attribution line]
[Primary CTA button] [Secondary: How?]
```

7.2 Per-Persona Hero Cards

PM — Meera

Headline: Your sprint has 2 hidden blockers — here's what to say before your standup

Temporal Hook: "Before your standup"

Stat Pills:

- `2 unresolved blockers` (Verified — Jira)
- `1 at-risk ticket` (Inferred — Jira + Slack)
- `Standup in 38 min` (Verified — Calendar)

Primary Insight: Two tickets marked In Progress have had no updates in 3 days and the assigned engineers haven't mentioned them in Slack. They are likely blocked but haven't been flagged. Artisan has drafted a check-in message for each.

Trust Source Attribution: From Jira sprint data + Slack channel analysis (last 72h)

Primary CTA: Review draft check-ins

Analyst — Arjun

Headline: 3 of your 5 Slack data requests this week were the same query

Temporal Hook: "Before engineering finds them"

Stat Pills:

- `3 duplicate requests` (Verified — Slack)
- `42 min/week recoverable` (Inferred — activity pattern)
- `1 auto-reply draft ready` (Suggested — Artisan)

Primary Insight: The weekly active users query appears in 3 different Slack threads with minor variations. Artisan has built a self-serve response template. Next time the question appears, you can reply in one click.

Trust Source Attribution: From Slack thread analysis (last 30 days)

Primary CTA: Review auto-reply template

Marketing Manager — Vikram

Headline: Your email campaign outperformed paid by 3.2x — here's why it won't last

Temporal Hook: "Before your Monday sync"

Stat Pills:

- `3.2x email vs. paid CTR` (Verified — Mailchimp + GA)
- `Open rate plateauing` (Inferred — trend analysis)
- `Budget reallocation window: this week` (Suggested — Artisan)

Primary Insight: Email is outperforming paid search this week, but open rates have flattened for 3 consecutive sends — a leading indicator of list fatigue. The window to reallocate budget to capitalize on email's current edge while refreshing the list is this week.

Trust Source Attribution: From Mailchimp campaign data + Google Analytics (last 21 days)

Primary CTA: See reallocation recommendation

Designer — Vagisha

Headline: Your last 3 design reviews averaged 47 minutes of re-explaining context

Temporal Hook: "Before your next design crit"

Stat Pills:

- `47 min avg context re-explain` (Inferred — meeting transcripts)
- `62% reviewers arrived unprepared` (Inferred — Figma open data)
- `Next crit: Thursday 2pm` (Verified — Calendar)

Primary Insight: Across your last 3 design crits, an average of 47 minutes was spent re-explaining constraints and prior decisions that were already documented in Figma or Notion. Artisan can generate a pre-crit brief for Thursday's session that reviewers receive 24 hours in advance.

Trust Source Attribution: From Calendar + Figma activity logs + meeting transcripts

Primary CTA: Generate Thursday's pre-crit brief

Sales — Shubham

Headline: Your 2pm call — here's what Infosys cares about right now

Temporal Hook: "5 minutes before your call"

Stat Pills:

- `Last contact: 12 days ago` (Verified — CRM)
- `3 open questions from last call` (Verified — Gong)
- `Recent news: Q3 hiring freeze` (Verified — web signal)

Primary Insight: Your last call ended with three unresolved questions about integration complexity. The prospect's LinkedIn shows a recent post about Q3 budget constraints — the hiring freeze context is directly relevant to your ROI framing. Artisan has drafted an opening that addresses both.

Trust Source Attribution: From Salesforce + Gong transcript + LinkedIn signal

Primary CTA: Review call brief + opening line

8. Artisan as Trained Subordinate — Vagisha Trust-Building Bot

The Trust-Building Bot represents Artisan's most sophisticated capability demonstration: not just surfacing information, but proving to a skeptical power user (Vagisha) that Artisan can function as a prepared, reliable subordinate — one who does the work before being asked and can show their work when questioned.

8.1 The Design Review Workflow (8 Steps)

The design review workflow is the primary use case for Vagisha and follows this sequence:

```
Step 1: Design file published in Figma
  ↓
Step 2: Artisan detects new file / major revision via Figma API
  ↓
Step 3: Artisan generates design brief (constraints, goals, 3 decisions needed)
  ↓
Step 4: Brief sent to invited reviewers 24h before crit
  ↓
Step 5: Recording Bot joins design crit call
  ↓
Step 6: Real-time transcription + decision tagging during call
  ↓
Step 7: Post-call: ReviewTrace auto-populated
  ↓
Step 8: Decision log shared with team; Figma comments updated with attribution
```

8.2 The Trust-Building Flow

The trust-building flow is the sequence of interactions that moves Vagisha from skeptic to advocate. It follows five phases:

Phase 1 — Capability Announcement

Artisan does not pitch itself. It announces a specific, concrete capability in the context of work already happening: "I noticed you have a design crit on Thursday. I can prepare a brief for your reviewers. Want to see what it would look like?"

Phase 2 — Evidence Surface

Before Vagisha asks "can it do this?", Artisan demonstrates it can — by generating a draft brief from the actual Figma file, with the actual constraints, and the actual three questions that need decisions. The evidence is specific to this design, not generic.

Phase 3 — Pre-Answered Expert Questions

Artisan anticipates the questions a skeptical expert would ask about its output and answers them proactively:

- "How did you determine these are the three decisions that need input?" → Source trace shown
- "What if a reviewer hasn't seen the prior context?" → Brief includes prior context summary
- "How do you handle conflicting feedback?" → ReviewTrace flags conflicts; Vagisha resolves

Phase 4 — Trace Availability

Every output Artisan produces in the design workflow is traceable. Vagisha can tap any element of the ReviewTrace and see exactly which source (transcript timestamp, Figma comment, Notion doc section) produced it. Nothing is asserted without evidence.

Phase 5 — Calibration Loop

After the first review cycle, Artisan surfaces a calibration prompt: "The brief I generated for Thursday's crit — was it accurate? Were the three questions the right ones?" Vagisha's response directly adjusts the brief generation model for future reviews.

8.3 The ReviewTrace Data Model

```
interface ReviewTrace {
  reviewId: string;
  designFileId: string;           // Figma file ID
  designFileVersion: string;      // Figma version at time of review
  reviewDate: ISO8601;

  decisions: Decision[];
  openQuestions: Question[];
  conflicts: Conflict[];
  participants: Participant[];
}

interface Decision {
  decisionId: string;
  summary: string;                // 1-2 sentence description
  madeBy: Participant;           // attributed reviewer
  madeAt: ISO8601;
  sourceType: 'transcript' | 'figma_comment' | 'slack' | 'notion';
  sourceRef: string;             // timestamp, comment ID, or doc section
  trustMarker: 'verified' | 'inferred';
  relatedDesignComponent?: string; // Figma node ID if applicable
  conflictsWith?: string[];       // decisionIds that conflict
}

interface Question {
  questionId: string;
  text: string;
  askedBy: Participant;
  status: 'open' | 'resolved' | 'deferred';
  resolution?: string;
  resolvedAt?: ISO8601;
}

interface Conflict {
  conflictId: string;
  decisionIds: string[];         // the two (or more) conflicting decisions
}
```

```

    description: string;
    resolvedBy?: Participant;
    resolution?: string;
  }

  interface Participant {
    participantId: string;
    name: string;
    role: string;
    attended: boolean;
    briefOpened: boolean;           // did they open the pre-crit brief?
    briefOpenedAt?: ISO8601;
  }

```

8.4 The Two-Sided Assessment Architecture

The design review pipeline uses two specialized bots:

Recording Bot

Role: Passive observer during the design crit call

Capabilities:

- Joins calendar-linked call (Google Meet / Zoom) automatically
- Performs real-time transcription
- Tags utterances with speaker identification
- Flags potential decision moments (phrases like "let's go with", "we've decided", "I think we should") for human review
- Does not intervene in the meeting

Output: Raw transcript + decision-candidate timestamp log

Assessment Bot

Role: Post-call processor and ReviewTrace generator

Capabilities:

- Ingests Recording Bot output
- Cross-references with Figma file, Notion brief, and prior ReviewTraces for the same project
- Classifies flagged moments as: Decision / Question / Action Item / Context
- Resolves ambiguous attributions using speaker ID + conversation context
- Generates the ReviewTrace
- Flags conflicts with prior decisions
- Drafts Figma comment updates with decision attribution

Output: Completed ReviewTrace + draft Figma annotations + conflict flags for Vagisha review

The two-sided architecture keeps the meeting experience clean (Recording Bot is invisible) while ensuring the post-meeting output is high-quality and fully traced (Assessment Bot does the heavy lifting after the fact).

9. Data Schema

9.1 Behavioral Events

All user interactions with Artisan generate structured behavioral events that feed into the model training pipeline.

```
// Core event schema
interface ArtisanEvent {
  eventId: string;
  userId: string;
  sessionId: string;
  timestamp: ISO8601;
  eventType: EventType;
  payload: Record<string, unknown>;
  source: 'user' | 'system' | 'integration';
}

type EventType =
  // Onboarding
  | 'onboarding.stage1.tool_connected'
  | 'onboarding.stage2.question_answered'
  | 'onboarding.stage3.hero_card_shown'
  | 'onboarding.stage3.hero_card_cta_clicked'
  | 'onboarding.stage3.hero_card_dismissed'
  | 'onboarding.stage3.hero_card_feedback'

  // Ritual engagement
  | 'ritual.morning_read.opened'
  | 'ritual.morning_read.item_acted_on'
  | 'ritual.pre_meeting_brief.opened'
  | 'ritual.pre_meeting_brief.opened_late' // opened after meeting started
  | 'ritual.capture.reviewed'
  | 'ritual.capture.item_edited'
  | 'ritual.eod_close.opened'
  | 'ritual.eod_close.dismissed'

  // Trust marker interactions
  | 'trust.marker.tapped'
  | 'trust.marker.source_viewed'

  // Card interactions
  | 'card.cta.primary_clicked'
  | 'card.cta.secondary_clicked'
  | 'card.feedback.positive'
  | 'card.feedback.negative'
  | 'card.feedback.corrected' // user corrects an inference

  // Integration events
  | 'integration.tool.connected'
  | 'integration.tool.disconnected'
  | 'integration.data.synced'
  | 'integration.anomaly.detected'
```

9.2 Computed Signals

Four key computed signals are derived from behavioral events and integration data. These signals are the primary inputs to Artisan's personalization and prioritization models.

Context Depth Score (CDS)

Measures how well Artisan knows the user — how many tools are connected, how much data has been processed, and how many training signals have been received.

```
CDS = (tools_connected × 0.3)
      + (days_of_data × 0.2)
      + (feedback_events × 0.3)
      + (correction_events × 0.2)
```

Range: 0-100

D7 target: >40

D30 target: >70

Morning Ritual Score (MRS)

Measures consistency of Morning Read engagement as a leading indicator of habit formation.

```
MRS = (morning_reads_opened / morning_reads_delivered) × 100
      × recency_weight // more recent opens weighted higher
```

Range: 0-100

Target: >60 at D14

Churn risk signal: MRS < 20 at D7

Shadow Tool Score (STS)

Measures whether Artisan is replacing the user's previous tool-checking behavior (the "shadow tools" they were using before Artisan).

```
STS = 1 - (shadow_tool_opens_post_artisan / shadow_tool_opens_pre_artisan)
```

```
Range: 0-1 (higher = more displacement)
Healthy signal: STS > 0.3 at D30
Computed from: browser activity, mobile app usage (where available)
```

Habit Streak

A simple consecutive-day engagement counter, surfaced in the EOD Close ritual as a variable reward.

```
habit_streak = consecutive_days where:
  morning_read.opened OR
  pre_meeting_brief.opened OR
  eod_close.opened
```

9.3 ML Training Signals

Signal	Description	Used For
`hero_card.cta.clicked`	User clicked primary CTA on hero card	Card relevance model
`hero_card.feedback.negative`	User marked card as not relevant	Persona model recalibration
`card.feedback.corrected`	User corrected an inference	Trust marker accuracy improvement
`ritual.morning_read.item_acted_on`	User acted on a Morning Read item	Item prioritization model
`ritual.pre_meeting_brief.opened_late`	Brief opened after meeting started	Notification timing model
`integration.anomaly.detected` + follow-up action	User acted on anomaly flag	Anomaly relevance model
`trust.marker.source_viewed`	User inspected a trust source	Trust transparency demand signal

10. Integration Layer


10.1 Dynamo V2 / Clarity MCP Server Mapping

Artisan's integration layer is built on the Clarity MCP Server, which abstracts tool-specific APIs into a unified capability surface. The mapping below shows which Clarity MCP tools power which Artisan capabilities.

Artisan Capability	Clarity MCP Tool	Underlying Integration	Data Type
Sprint blocker detection	`clarity.jira.sprint_health`	Jira REST API	Ticket status, assignee activity
Stakeholder update draft	`clarity.confluence.doc_draft`	Confluence API	Page templates, sprint data
Slack blocker inference	`clarity.slack.thread_analysis`	Slack Events API	Thread activity, mention patterns
Campaign performance brief	`clarity.analytics.aggregation`	GA4 + HubSpot API	Campaign metrics, delta computation
Pre-meeting brief	`clarity.calendar.meeting_context`	Google Calendar API	Event metadata, attendee history
Call prep card	`clarity.crm.account_context`	Salesforce / HubSpot API	Account history, open opportunities
Design crit brief	`clarity.figma.file_summary`	Figma REST API	File version, component changes
ReviewTrace generation	`clarity.transcript.decision_extract`	Meet / Zoom transcript API	Speaker-attributed transcript
Anomaly detection	`clarity.signals.anomaly_detect`	Any connected integration	Time-series deviation
CRM draft update	`clarity.crm.activity_draft`	Salesforce / HubSpot API	Call transcript → CRM fields

10.2 MCP Tunnel Architecture for Enterprise

Enterprise customers require that no raw data leaves their network. Artisan supports a local MCP tunnel architecture:



Tunnel Agent responsibilities:

- Authenticates to on-premise tools using local credentials

- Extracts structured signals (ticket counts, thread metadata, calendar events) without forwarding raw content
- Applies org-defined data classification rules before transmission
- Caches signals locally for 24h to support offline operation

What never leaves the tunnel:

- Email body content
- Document full text
- Call recordings or raw transcripts
- PII beyond what is required for attribution (name + role)

What the tunnel sends:

- Aggregated metrics (ticket counts, velocity, response times)
 - Metadata (titles, statuses, timestamps, assignees)
 - Structured decision outputs from on-premise Assessment Bot
 - Behavioral event signals (anonymized by default)
-

11. Open-Source Integration Map

Artisan is designed with a first-class integration path for open-source alternatives to the commercial SaaS tools most knowledge workers use. The three anchor open-source integrations are: **Twenty CRM**, **AppFlowy**, and **Cal.com**.

11.1 Integration Map

Dimension	Twenty CRM	AppFlowy	Cal.com
Integration Layer	REST API + webhook	REST API + real-time sync	REST API + webhook
Tool Discovery Signal	Self-hosted CRM presence on network scan; Twenty-specific API response headers	AppFlowy workspace metadata; self-hosted document structure	Cal.com booking page pattern; self-hosted calendar API endpoint
Interview Stage Pre-fill	Pull open opportunities, recent contact activity, pipeline stage	Pull recent documents, workspace structure, tag taxonomy	Pull calendar density, recurring meeting patterns, availability windows
Personalization Card	"Your pipeline has X deals with no activity in 14 days"	"Your most-accessed workspace is [X] — Artisan can brief you before each session"	"You have X discovery calls this week — briefs ready for all of them"
Daily Brief Delta	Deal activity changes since last brief; new contacts added	Document updates in watched workspaces; new collaborative edits	Upcoming meetings for next 24h; new bookings since last brief
Task Lifecycle Source	Opportunity stages → task triggers (e.g., deal moves to Proposal → trigger follow-up brief)	Document status changes → brief triggers	Meeting booked → automatic pre-meeting brief trigger
Trust Marker	Verified: direct API read from Twenty CRM instance	Verified: direct API read from AppFlowy instance	Verified: direct API read from Cal.com instance
MCP Opportunity	`clarity.twenty.pipeline_health` — wraps Twenty CRM API for pipeline signal extraction	`clarity.appflowy.workspace_summary` — wraps AppFlowy API for document signal extraction	`clarity.calcom.meeting_context` — wraps Cal.com API for calendar signal extraction
Self-Hosted Advantage	Full data residency; no HubSpot/Salesforce licensing cost; API without rate limits	Full document privacy; no Notion data sharing; local vector search on documents	Full calendar privacy; no Google Calendar data sharing; custom booking workflows
Persona Best Fit	Sales (Shubham) — primary CRM workflow; also PM for stakeholder tracking	PM (Meera), Designer (Vagisha), Analyst (Arjun) — document-heavy workflows	All personas — universal calendar integration; especially Sales for booking-driven workflows

11.2 Self-Hosted Stack Configuration

For users running the full open-source stack, Artisan's recommended self-hosted configuration is:

```
# artisan-self-hosted.yml
integrations:
  crm:
    provider: twenty
    endpoint: https://crm.internal.company.com
    auth: api_key

  documents:
    provider: appflowy
    endpoint: https://docs.internal.company.com
    auth: oauth2

  calendar:
    provider: calcom
    endpoint: https://cal.internal.company.com
    auth: api_key

mcp_tunnel:
  enabled: true
  data_residency: local
  transmission_mode: structured_signals_only

artisan_cloud:
  endpoint: https://api.artisan.ai
  auth: org_token
  features:
    personalization: true
    card_generation: true
    anomaly_detection: true
```

12. Success Metrics

Success for Artisan is measured across three horizons: onboarding (Day 1), early retention (Day 7), and habit formation (Day 30).

12.1 Onboarding Metrics (Day 1)

Metric	Definition	Target
Stage 1 Completion Rate	% of users who connect at least 1 tool	>85%
Stage 2 Completion Rate	% of users who complete the Interview stage	>70%
Stage 3 Completion Rate	% of users who reach the hero card	>65%
Onboarding Duration	Median time from signup to hero card display	<8 minutes
Hero Card CTA Click Rate	% of users who click the primary CTA on the hero card	>40%
Wow Moment Conversion Rate	% of users who click CTA AND return within 24h	>30%

12.2 Early Retention Metrics (Day 7)

Metric	Definition	Target
D7 Retention	% of users active on Day 7	>45%
Context Depth Score at D7	Median CDS across active users at Day 7	>40
Morning Ritual Score at D7	% of users with MRS > 60 at Day 7	>25%
Capability Card CTR	% of capability cards that result in a primary CTA click	>35%
Trust Marker Interaction Rate	% of users who tap at least one trust marker in first 7 days	>20%

12.3 Habit Formation Metrics (Day 30)

Metric	Definition	Target
D30 Retention	% of users active on Day 30	>30%
Context Depth Score at D30	Median CDS across active users at Day 30	>70
Shadow Tool Score at D30	% of users with STS > 0.3 at Day 30	>35%
Morning Ritual Score at D30	% of users with MRS > 60 at Day 30	>40%

Metric	Definition	Target
Habit Streak at D30	Median habit streak for retained users at Day 30	>12 days
Stage 3 Wow Moment to D30 Retention	% of users who clicked hero card CTA and are retained at D30	>55%

12.4 Quality Metrics

Metric	Definition	Target
Inference Accuracy	% of role inferences confirmed correct by users	>80%
Card Correction Rate	% of cards where user corrects an inference	<15%
Trust Marker Dispute Rate	% of trust markers disputed (user says source is wrong)	<5%
Anomaly False Positive Rate	% of anomaly flags dismissed as not relevant	<25%
ReviewTrace Completeness	% of decisions in a design crit captured in ReviewTrace	>80%

12.5 Metric Hierarchy

Primary (board-level):

- D30 Retention → target: 30%
- Wow Moment Conversion Rate → target: 30%

Secondary (growth-level):

- Stage 3 Completion Rate
- Hero Card CTA CTR
- Context Depth Score D7

Tertiary (product-level):

- Morning Ritual Score
- Shadow Tool Score
- Inference Accuracy
- ReviewTrace Completeness

13. Reliance Score — The Right Company Metric

13.1 The Problem with Engagement Metrics

Engagement metrics — DAU, session length, feature usage rate — measure activity, not value. A user who opens Artisan every morning and ignores the suggestions is "engaged." A user who delegates two critical workflows and never looks back is irreplaceable. These are different things.

The goal is not to be used. The goal is to be irreplaceable.

This distinction changes everything about product design:

- Engagement-optimized products create habit loops (notifications, streaks, friction removal)
- Irreplaceability-optimized products deepen context, improve calibration, and expand the scope of what users dare to delegate

13.2 Reliance Score Definition

Reliance Score (RS) measures how much of a user's meaningful work Artisan has become irreplaceable for.

RS is computed daily, per user, as a weighted composite:

$$RS = (\text{context} \times 0.30) + (\text{calibration} \times 0.30) + (\text{proactive} \times 0.25) + (\text{orchestration} \times 0.15)$$

Dimension	Weight	What It Measures	Example Signal
Context Depth	30%	How much of the user's world Artisan knows	Tools connected, projects tracked, decisions remembered
Calibration	30%	How much Artisan has learned the user's taste and corrected mistakes	Correction rate declining, user accepting suggestions increasing
Proactive Value	25%	% of user's proactive sessions where Artisan surfaced something they acted on	Opened morning brief → clicked through → took action
Orchestration	15%	User has delegated multi-step workflows Artisan now owns	Review bot runs, standup generation auto-sent, report auto-filed

RS is a **0–1 score**. The company goal is:

> **Move the RS floor from 0.44 to 0.70 for L5 skilled users within 6 months of account creation.**

The floor matters more than the average. If the median L5 user scores 0.72 but 40% of L5 users score 0.44, Artisan is not irreplaceable — it's optional.

13.3 The Dhanush Test

The Dhanush Test is the internal calibration for measuring what L1–L5 means:

Level	Description	Typical RS
L1	Opens Artisan, reads cards, takes no action	0.05–0.15
L2	Clicks CTAs, tries a workflow once	0.15–0.30
L3	Uses 2–3 workflows regularly, still has alternative path	0.30–0.50
L4	Artisan is the primary path for 1 high-stakes workflow	0.50–0.65
L5	Cannot imagine removing Artisan from daily flow	0.65–0.90

L5 is defined behaviorally, not by feature usage. The signal is: **would their day break without it?**

Named internally after Dhanush, the first internal user who reached L5 without being prompted to.

13.4 RS and Onboarding Design Implications

RS score has direct implications for how onboarding should be designed:

The goal of onboarding is not to teach features. It is to cross the L2 → L3 threshold before the user leaves.

This means:

- **Stage 1 (Connect)**:** Build context. Every integration connected increases RS potential ceiling.
- **Stage 2 (Interview)**:** Calibrate. Every answer reduces the number of corrections Artisan will need to make.
- **Stage 3 (Personalization)**:** Trigger first proactive value. The hero card should be specific enough that the user says "wait, how did it know that?"

If a user exits Stage 3 without feeling recognized, they will not reach L3.

13.5 RS as the Board Metric

The north star for Artisan is:

> "% of knowledge workers for whom removing Artisan would break their workflow (L5), measured at Day 90 post-signup"

Target: **25% of active users reach L5 by D90**

Supporting metrics:

- % reaching L3 by D7: target 60%
 - RS velocity (how fast RS rises in first 14 days): target 0.05/day
 - RS floor for cohort by D30: target 0.44
 - RS floor for cohort by D90: target 0.70
-

14. Oracle Open Cloud POC Brief

14.1 Background

ClearTax has an enterprise relationship with Oracle Cloud. The Artisan Oracle POC is a B2B demonstration: **can Artisan, connected to Oracle Open Cloud's MCP infrastructure, deliver L3+ RS for enterprise knowledge workers in under 90 minutes of onboarding?**

This is a B2B track running in parallel to the B2C/PLG track. It informs whether Artisan can be sold as an enterprise AI work surface vs. a consumer product.

14.2 POC Structure

Duration: 2 weeks

Participants: 8 Oracle Cloud users across 3 roles (PM × 3, Analyst × 3, EM × 2)

Tools connected: Oracle HCM, Oracle Analytics Cloud, Oracle JIRA-equivalent (Project Management), Slack, Google Workspace

Success criterion: 5 of 8 users reach RS ≥ 0.50 by Day 14

14.3 Five POC Tests

Test	What We're Testing	Measurement
T1 — Onboarding Speed	Can an enterprise user complete all 3 stages in under 12 minutes?	Time to Stage 3 completion
T2 — Context Quality	Does Artisan correctly infer role and stack from Oracle Cloud data alone?	Inference accuracy confirmed by user
T3 — Hero Card Resonance	Does the hero card reference something real about the user's current work?	"Yes, this is about my current project" rate
T4 — First Delegation	Does the user voluntarily initiate a second workflow within 48 hours?	% of users returning within 48h
T5 — Trust Transfer	Does the user share an Artisan output with a colleague unasked?	Forward / share rate

14.4 Hypotheses

H1 (Onboarding): Enterprise users are more willing to connect integrations because they expect data to be shared across internal tools. They will connect more integrations than consumer users.

H2 (Context): Oracle Analytics Cloud data will produce higher-quality hero cards than Google Workspace alone because it contains structured metrics and project data.

H3 (Trust): Enterprise users will trust Artisan outputs more if the trust source references an internal system (e.g., "Source: Oracle HCM sprint data") vs. a consumer tool.

H4 (Delegation): Enterprise users will delegate workflows faster if Artisan pre-answers the question their manager would ask ("Amit asks Apur" pattern) — reducing the user's political risk in sharing AI output.

14.5 MCP Tunnel Architecture for Oracle

The POC uses the Clarity MCP Server as an enterprise tunnel:

```
Oracle Cloud APIs
  ↓
Clarity MCP Server (runs inside Oracle network, no public exposure)
  ↓ (SSE/HTTP JSON-RPC)
Artisan Frontend ← VITE_CLARITY_BACKEND_URL points to tunnel endpoint
  ↓
User sees oracle-native hero cards
```

The tunnel architecture means Oracle's data never leaves their network. Artisan calls the tunnel; the tunnel calls Oracle APIs; results are returned. This is the same pattern as the ClearTax deployment but with Oracle Cloud as the MCP host.

14.6 POC Success and Decision Gate

If $T3 \geq 75\%$ and $T4 \geq 50\%$, we proceed to a 90-day enterprise deployment with Oracle.

If $T3 < 50\%$, we redesign the hero card generation logic for structured Oracle data.

If $T4 < 30\%$, we investigate whether the "first delegation" trigger needs to be more explicit in Stage 3.

15. The Four Onboarding States

15.1 Why Four States, Not a Funnel

Traditional onboarding funnels measure completion rates: % who completed Step 1, Step 2, Step 3. This treats onboarding as a sequence of tasks. Artisan's onboarding is a sequence of **belief changes**. Each state is a new mental model the user has about what Artisan is for.

The four states:

- **Recognition** — "It sees my world"
- **Discovery** — "I didn't know this was possible"
- **Belief Change** — "My mental model shifted"
- **Trust + Repeat** — "I will delegate again"

A user can exit onboarding (click through Stage 3) without ever reaching State 2. That user will not return. The onboarding design goal is to **trigger all four states within the onboarding session, not after it**.

15.2 State 1: Recognition — "It sees my world"

Trigger moment: End of Stage 1 (Connect), when the scanning animation resolves and shows discovered tools.

What must happen for Recognition to occur:

- Artisan must name at least one real tool the user uses
- The tool must be named specifically ("Linear" not "project management tool")
- The display must feel personal, not generic

The Loading Experience (90-second hook):

The scan animation is the highest-leverage moment in the entire onboarding. This is where Recognition either happens or doesn't.

Design requirements:

- Stage 1 scan must feel alive and intelligent, not like a progress bar
- Each integration scan must surface tool names as they are discovered: "Found Linear... Found Notion... Found Figma..."
- The final "scan complete" moment must show a count: "17 tools found across your workspace"
- The persona inference must appear as a consequence of the scan: "Based on your tools and activity, you look like a **Product Manager**"

What breaks Recognition:

- Generic output ("Connected to Google Workspace")
- Wrong tool names (hallucinated or out of date)
- Tool names the user doesn't recognize (wrong scan results)
- Persona inference that feels random

Recognition as the trust foundation:

Recognition is the only state that builds purely on *data accuracy*. It does not require AI output quality. If Artisan names real tools, users trust that Artisan sees something real. This is Stage 1 of the three trust angles (see 15.5).

15.3 State 2: Discovery — "I didn't know this was possible"

Trigger moment: Stage 3, when the hero card is first shown.

What must happen for Discovery to occur:

- The hero card must reference something specific about the user's current reality (not generic)
- The CTA must describe a capability the user didn't know existed
- The "wow, Artisan can do *that*?" reaction must occur

Discovery triggers per persona:

Persona	Discovery Hook	What the user didn't know was possible
PM (Meera)	"Sprint ends in 3 days. 2 tickets haven't moved."	Artisan can proactively surface stuck tickets before standup
Analyst (Arjun)	"Your weekly metrics report takes 90 minutes. Here's the same report in 4."	Artisan can generate the exact report format Arjun uses
Marketing (Vikram)	"HubSpot says 12%. GA says 8.3%. One of them is wrong."	Artisan can reconcile attribution conflicts across tools
Designer (Vagisha)	"Handoff is in 3 days. The spec isn't written yet."	Artisan can write the spec from Figma annotations
Sales (Shubham)	"7 of your deals haven't been touched in 21 days."	Artisan can generate outreach drafts ranked by close probability

Discovery is not about features. It is about reframing what the user thought their job was.

A PM who discovers that Artisan can write the standup brief now has a 10-minute window back every morning. The discovery isn't "Artisan writes standups." The discovery is "I don't have to do that anymore."

15.4 State 3: Belief Change — "My mental model shifted"

Belief Change is the hardest state to engineer because it cannot be forced. It is the result of accumulated small surprises.

The Belief Map (B0 → B5):

Belief State	What the user believes	What triggers the move
B0	"AI can do basic things but not my real work"	N/A (entry state)
B1	"Artisan knows my tools"	Recognition (State 1)
B2	"Artisan can surface something I missed"	First accurate hero card (State 2)
B3	"Artisan can draft something I would have written"	First accepted output (not just read — accepted)
B4	"I can review Artisan's work instead of doing it myself"	First delegation where the user modified output rather than rewrote from scratch
B5	"I am a full-stack builder / PM / analyst now"	First time user ships something they couldn't have shipped before Artisan

Important: B states are sequential. Users cannot skip from B1 to B4. Product design must create the conditions for B2 to happen before attempting to deliver B3-level outputs.

Belief Change in onboarding: The goal of Stage 3 is to take the user from B0/B1 to B2 during the onboarding session itself. B3+ happens in the product's main surface (the Artisan workspace). If the user leaves Stage 3 at B1 or below, they will not reach B3 organically.

Belief Change and the full-stack designer:

The most radical belief change is for the Designer persona. A UI designer today:

- Designs a component in Figma
- Writes a spec
- Hands off to engineering
- Engineering makes changes
- Designer reviews and approves

Tomorrow, with Artisan at B5:

- Designs a component in Figma
- Artisan writes the spec from annotations
- Artisan generates the component code (React/Tailwind/shadcn)
- Designer reviews the code output, not the spec
- Designer approves and ships to production

The Designer is no longer a designer who hands off to engineering. She is a **full-stack builder** who also does design. This is the B5 belief shift for the designer persona. Artisan's job is to create the conditions where this shift feels natural, not forced.

15.5 Three Trust Angles

Trust in Artisan comes from three independent sources. Each must be satisfied for the user to reach B4+:

Trust Angle 1: Data Trust (Does Artisan see the right things?)

- Built by: Recognition (State 1), accurate tool names, real project names
- Broken by: Wrong tool names, hallucinated project references, stale data
- Design principle: Never show a data point you can't trace. The Trust Marker system exists for this.

Trust Angle 2: Emotional Trust (Does Artisan get my context?)

- Built by: Tone matching (card language that sounds like the user's domain), persona inference accuracy, temporal hooks ("before your standup" not "in the morning")
- Broken by: Generic language, wrong persona, outputs that feel like they were written for someone else
- Design principle: Every card must contain at least one phrase that could not have been written for a different user.

Trust Angle 3: Human-in-the-Loop Trust (Can I correct Artisan without losing value?)

- Built by: Easy correction mechanisms (persona corrector in Stage 3, card dismissal, output editing), calibration signals that improve over time, trace layer showing **why** Artisan decided something
- Broken by: Outputs that ignore corrections, trace layer that is too complex to act on, correction that resets all learning
- Design principle: Correction must be a signal, not a punishment. When a user corrects Artisan, RS goes up (calibration dimension), not down.

15.6 State 4: Trust + Repeat — "I will delegate again"

Trigger moment: The user voluntarily initiates a second Artisan workflow within 48 hours of onboarding completion.

This is the "aha" moment in the behavioral sense. It is the moment when Artisan stops being a product the user tried and becomes a habit.

What makes State 4 happen:

- The first workflow must have produced something the user actually used
- The first workflow must have been faster than the user's existing process
- The user must have a second obvious workflow to delegate (Artisan surfaces this proactively)

The "next workflow" surface:

At the end of every Artisan output, the UI shows: "Want me to also handle X?" — where X is the next most likely workflow based on what was just completed. This is the State 4 trigger. It is not a feature suggestion. It is an invitation based on a specific observed context.

Example sequences:

- PM: "Standup brief generated → Want me to also update the sprint tracker with these blockers?"
- Analyst: "Weekly report generated → Want me to schedule this to auto-send every Monday at 9am?"

- Designer: "Spec written → Want me to generate the component code from this spec?"

15.7 Loading Experience Design

The loading experience across all three stages is critical for building the emotional arc. The user's mental model of Artisan is being formed in the 90 seconds between pressing "Connect" and seeing Stage 3.

Stage 1 Loading (Scan):

- Integration icon appears, then scans with a pulsing ring
- Tool names appear below the integration as they are "discovered"
- Counter increments: "14 tools found... 15 tools found..."
- Scan complete: Brief pause, then "17 tools across 4 integrations. Let's learn how you actually use them."

Stage 2 → Stage 3 Transition:

- After the last interview question is answered, show: "Analyzing your tools and answers..."
- Then: "Building your workspace profile..."
- Then: "Generating your personalized view..."
- Then: The hero card appears with a subtle fade-in (not a loading spinner → sudden appearance)

Timing principle: The Stage 1 scan should take 8–12 seconds per integration (even if the actual scan completes faster). This is "perceived intelligence time" — users trust results that took time more than results that appeared instantly. The delay signals processing, not performance issues.

16. The Amit Asks Apur Quality Report

16.1 The Pattern

In ClearTax's engineering culture, design reviews involve a recurring dynamic:

- **Amit** is the engineering manager. He reviews UI changes before merge. He has a checklist of concerns: accessibility, mobile responsiveness, design system compliance, legacy screen compatibility, ship risk.
- **Apur** is the designer/developer who submits the PR. He knows Amit will ask these questions. He could pre-answer them — but writing the answers into every PR body is tedious and inconsistent.

The Artisan intervention: Artisan generates the PR body automatically, pre-answering every question Amit will ask. Apur reviews and approves the generated body before submitting. Amit gets a PR that reads like Apur thought of everything.

The pattern generalizes: **any recurring output that must satisfy a known reviewer's checklist is a candidate for the "Amit asks Apur" pattern.**

16.2 Quality Report Structure

The auto-generated PR body follows this structure:

```
## What Changed
[1-3 sentences. What the PR does, not how.]

## Why
[1-2 sentences. What user need or product goal this addresses.]

## Coverage

| Check | Status | Notes |
|---|---|---|
| WCAG AA accessibility |  Pass | Contrast ratios verified. Button labels updated. |
| Mobile responsiveness |  Pass | Tested at 375px, 414px, 768px |
| Design system compliance |  Pass | Uses shadcn/ui primitives. No custom overrides. |
| Legacy screen compat |  Review | One legacy modal uses z-index 9999. Verified non-conflicting. |
| Ship risk |  Low | No API changes. Feature-flagged for 10% rollout. |

## Screenshots
[Auto-inserted before/after screenshots from Figma diff or Storybook]

## Test Plan
- [ ] Reviewed in Storybook (all states)
- [ ] Tested on Chrome, Safari, Firefox
- [ ] Reviewed against WCAG checklist
- [ ] PM sign-off: @[pm-handle]
```

```
## Artisan Trace
```

```
Source: Figma spec v2.3 (linked), Accessibility audit (linked), Mobile screenshot (auto-generated)  
Generated by Artisan on [timestamp]. Apur reviewed and approved.
```

16.3 The Trust Mechanics

The quality report achieves three trust outcomes simultaneously:

For Apur (submitter):

- Reduces the cognitive overhead of writing PR bodies
- Ensures nothing important is missed (the checklist is always complete)
- Shifts work from "write the PR body" to "review the PR body" — a lower-effort activity

For Amit (reviewer):

- PR body answers his questions before he asks them
- Reduces review time (he doesn't have to chase down accessibility or mobile answers)
- The "Artisan Trace" section tells him what sources Artisan used — he can verify if he wants, but usually doesn't need to

For the team:

- PRs become a knowledge artifact: the quality report captures *why* every decision was made
- New team members can read old PRs and understand design rationale
- The trace layer means Artisan's outputs are auditable, not black-box

16.4 Amit's Checklist (Full)

The full checklist Artisan pre-answers:

- **WCAG AA**: All interactive elements have accessible labels. Color contrast $\geq 4.5:1$ for normal text, $\geq 3:1$ for large text. Focus indicators visible.
- **Mobile**: Responsive at 320px, 375px, 414px, 768px. No horizontal scroll. Touch targets ≥ 44 px.
- **Design System**: All components from approved design system library. No inline styles that override design tokens. Typography uses approved scale.
- **Legacy Compatibility**: No z-index conflicts with existing modals. No CSS that breaks on IE11 (if applicable). No Web Components APIs not polyfilled.
- **Ship Risk**: API surface unchanged OR breaking change documented with migration path. Feature-flagged if $>5\%$ of users affected. Rollback plan exists.
- **Performance**: No new synchronous network calls in render path. Images optimized. Bundle size delta < 5 KB.
- **Security**: No user data logged to console. No external URLs hardcoded. No eval() or innerHTML with user data.

16.5 Generalization: The Pattern for Other Roles

Reviewer	Submitter	Context	What Artisan pre-answers
Amit (EM)	Apur (Designer)	UI PR review	Accessibility, mobile, design system, risk
Priya (VP Product)	Meera (PM)	Feature spec review	User research backing, metric hypothesis, rollback plan
CFO	Finance Analyst	Budget request	ROI model, assumptions table, downside scenario
Client	Account Manager	Proposal	Pricing rationale, timeline, dependency list, risk register
Board	Founder	OKR review	Metric sources, methodology, trend direction, action plan

The Artisan "Amit asks Apur" pattern is the generalized version of the Vagisha trust-building bot for non-design roles. The trust mechanics are identical: Artisan pre-answers the reviewer's checklist so the submitter looks like they thought of everything.

16.6 Implementation Notes

The quality report is generated from:

- Figma diff (via Clarity MCP → Figma integration)
- Git diff (via Clarity MCP → GitHub integration)
- WCAG audit (via Clarity MCP → accessibility-check tool)
- Mobile screenshot (via Clarity MCP → screenshot-service tool)
- Feature flag status (via Clarity MCP → LaunchDarkly/GrowthBook integration)

The `review-output` component renders the quality report in the Artisan workspace. The user can edit any field before submitting. The "Artisan Trace" section is always shown — it cannot be removed, only minimized. This is a trust-non-negotiable.

Document ends. Version 0.4. Sections 13–16 added: Reliance Score, Oracle POC, Four Onboarding States, Amit/Apur Quality Report.